

## Mini-Batch Gradient Descent (eg. $X=512,000$ , mini batch 512)

for  $t = 1, \dots, 1000$ :

Forward Prop on  $X^{(t)}$

$$\text{Compute cost: } J^{(t)} = \frac{1}{512} \sum_{i=1}^1 L(\hat{y}^{(i)}, y^{(i)}) + \frac{\lambda}{2 \times 512} \sum \|W\|_F^2$$


Backprop to compute gradient  $J^{(t)}$

$$W^{(t)} := W^{(t-1)} - \alpha dw^{(t)}, \quad b^{(t)} := b^{(t-1)} - \alpha db^{(t)}$$

Basically same as gradient descent but run on each smaller batch so as to save computation time

Batch: whole dataset  $(X^{(t)}, y^{(t)}) = (X, y)$

Stochastic: 1 training example  $(X^{(t)}, y^{(t)}) = (X^{(i)}, y^{(i)})$



Batch: too slow for each iterations


Stochastic: lose speeding from vectorization

Mini batch: Fastest learning, make progress without waiting to process entire dataset.

$$n = \{64, 128, 256, 512\}$$

$2^6 \quad 2^7 \quad 2^8 \quad 2^9 \rightarrow \text{CPU / GPU optimization}$

## Exponentially Weighted Average



$$V_t = \beta V_{t-1} + (1-\beta) \nabla_t$$

$\beta = 0.9 \approx$  looking at past 10 days.

$V_t$  is approximately  $\frac{1}{1-\beta}$

## Bias correction

Problem: initially  $V_t = 0$ , the line looks like the blue one.

Solution: Let  $V_t = \frac{V_t}{1-(\beta)^t}$

as  $t$  increase,  $1-(\beta)^t$  close to 1

## Gradient Descent with Momentum

On iteration  $t$ :

Compute  $dw, db$  --- on mini-batch

$$V_{dw} = \beta V_{dw} + (1-\beta) dw$$

$$V_{db} = \beta V_{db} + (1-\beta) db$$

$$w := w - \alpha V_{dw}, \quad b := b - \alpha V_{db}$$

not weight, bias,  
just features,  
could also be  $w_1, w_2, \dots$ .

## RMSProp

On iteration  $t$ :

compute  $dw, db$  --- on mini-batch

$$S_{dw} = \beta S_{dw} + (1-\beta) dw^2$$

$$S_{db} = \beta S_{db} + (1-\beta) db^2$$

$$w := w - \alpha \frac{dw}{\sqrt{S_{dw}}}, \quad b := b - \alpha \frac{db}{\sqrt{S_{db}}}$$

dampen out oscillation (can use larger  $\alpha$ )



slow down  
vertical and  
speed up horizontal

**Adam** : a combination of Momentum and RMSProp

$$V_{dw}, S_{dw}, V_{db}, S_{db} = 0$$

on iteration  $t$ :

compute  $dw, db$  use mini-batch

$$V_{dw} = \beta_1 V_{dw} + (1 - \beta_1) dw, V_{db} = \beta_1 V_{db} + (1 - \beta_1) db$$

$$S_{dw} = \beta_2 S_{dw} + (1 - \beta_2) dw^2, S_{db} = \beta_2 S_{db} + (1 - \beta_2) db^2$$

$$V_{dw}^{corrected} = V_{dw} / (1 - \beta_1^t), V_{db}^{corrected} = V_{db} / (1 - \beta_1^t)$$

$$S_{dw}^{corrected} = S_{dw} / (1 - \beta_2^t), S_{db}^{corrected} = S_{db} / (1 - \beta_2^t)$$

$$w = \alpha \frac{V_{dw}^{corrected}}{\sqrt{S_{dw}^{corrected}} + \epsilon}$$

$$b = \alpha \frac{V_{db}^{corrected}}{\sqrt{S_{db}^{corrected}} + \epsilon}$$

avoid division by zero, typically  $10^{-8}$

**hyperparameters**:  $\alpha \rightarrow \text{tune}$   $\beta_1 = 0.9$   $\beta_2 = 0.999$   $\epsilon = 10^{-8}$


**Learning Rate Decay** : slowly decrease your learning rate.

$$\alpha = \alpha_0 \frac{1}{1 + \text{decay rate} * \text{epoch number}}$$

exponential decay  $0.9^{\text{epoch-num}}$   $\alpha^0$

$$\alpha = \frac{k}{\sqrt{\text{epoch-num}}} \cdot \alpha_0$$

Discrete decay



Manual Decay